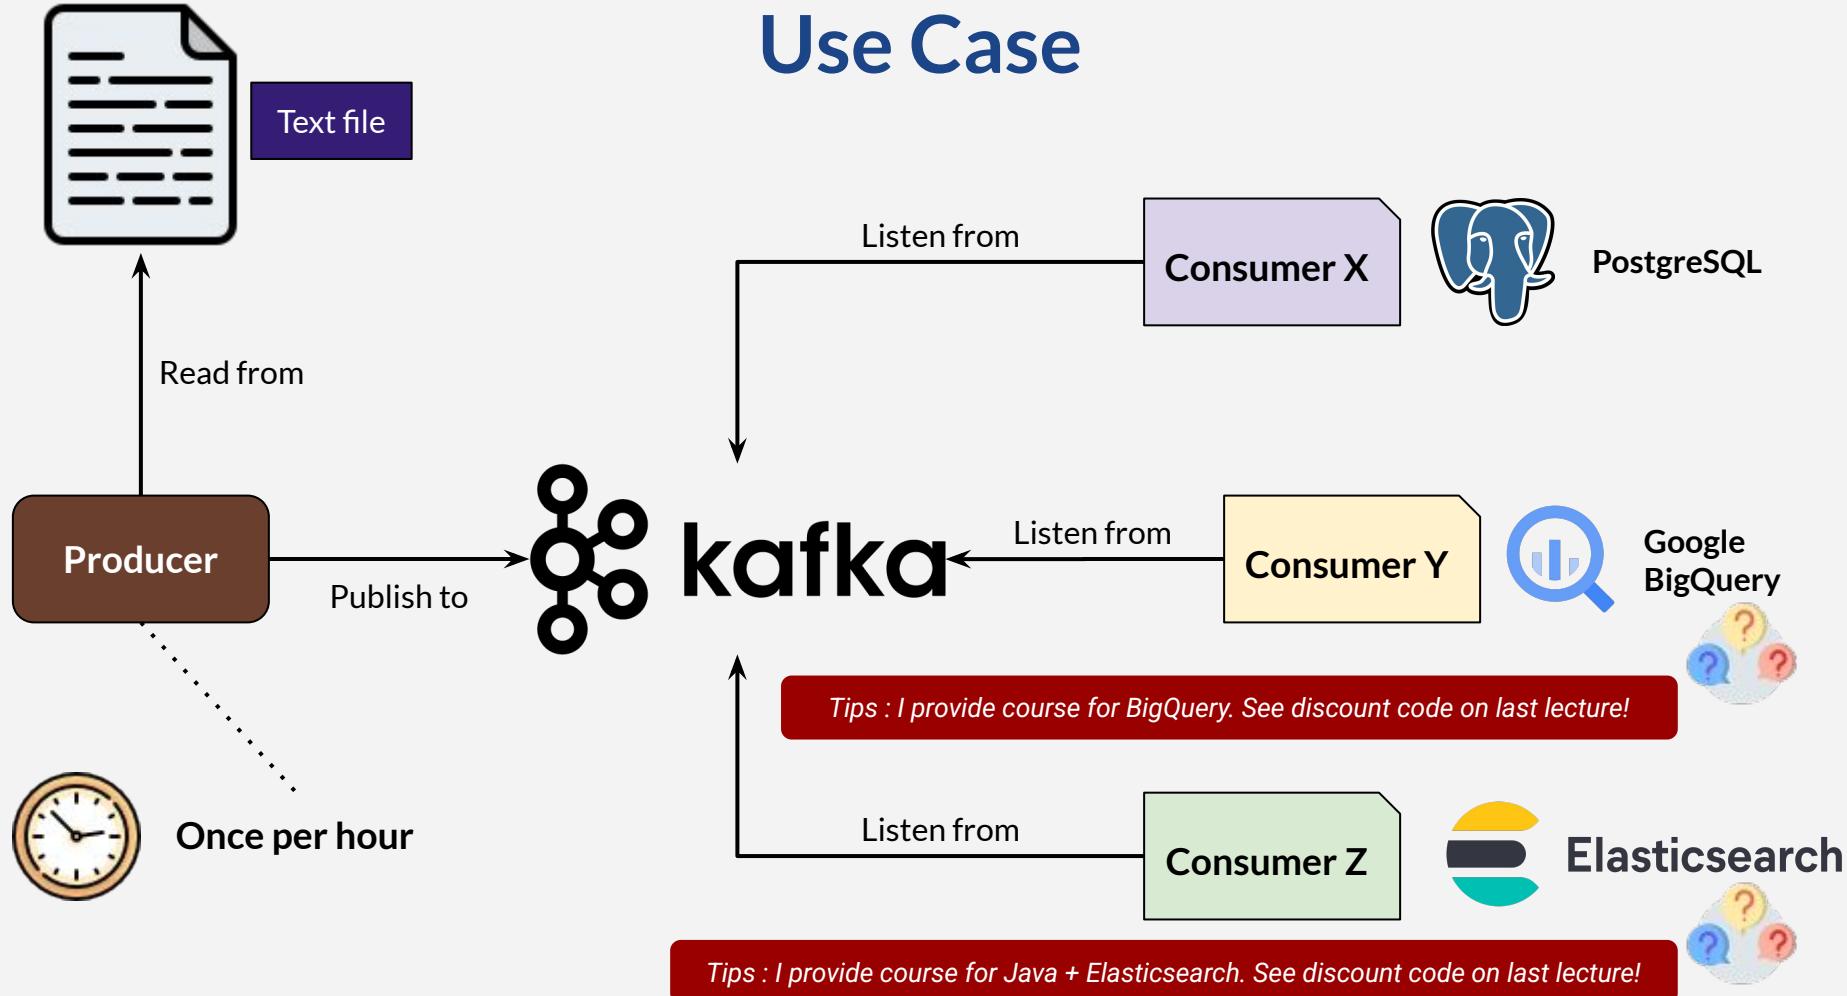


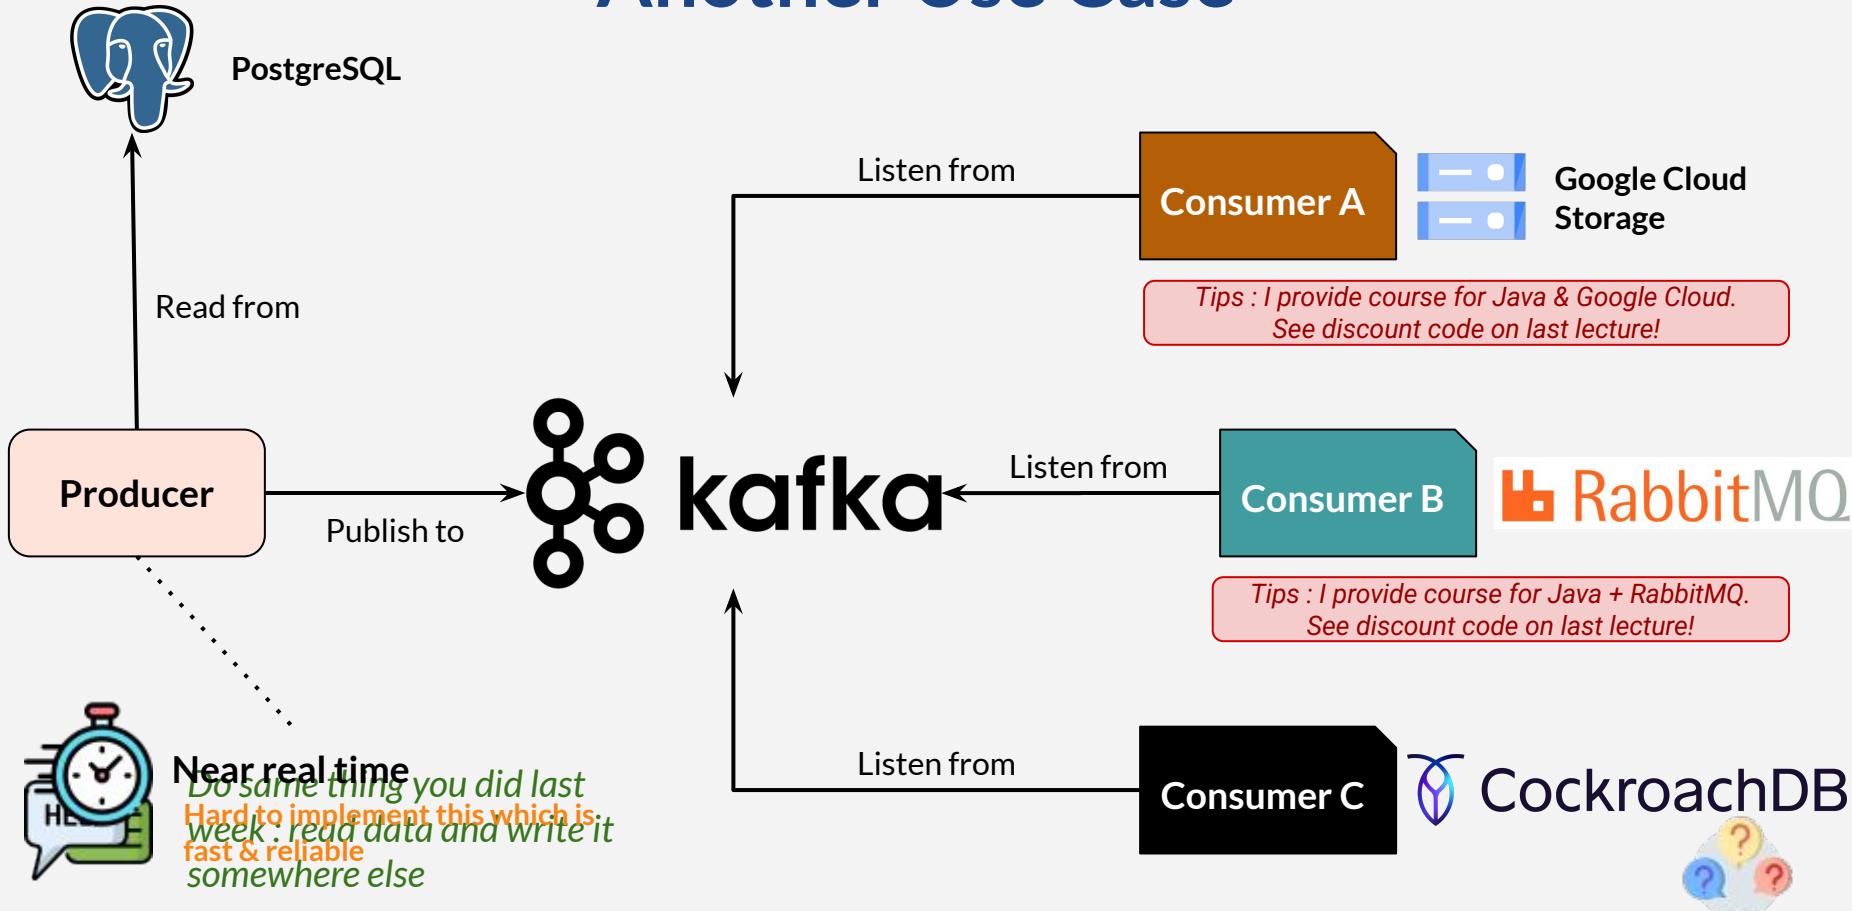
Message In, Message Out



Use Case



Another Use Case



Message In, Message Out

- × Happens in real life
- × Today we have multiple data sources (**in**)
 - × text file, relational databases
 - × Non relational database, email, cloud, social media, specific software, etc
 - × Can be more than 100
- × Multiple target data store (**out**)
 - × local text file, FTP, relational / non-relational databases, big data, cloud, etc



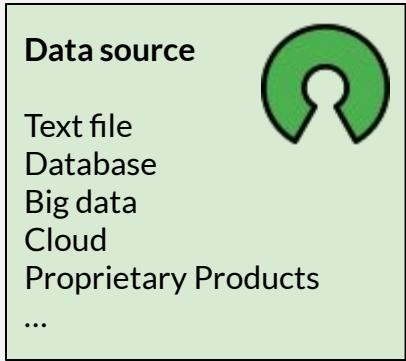
On Kafka

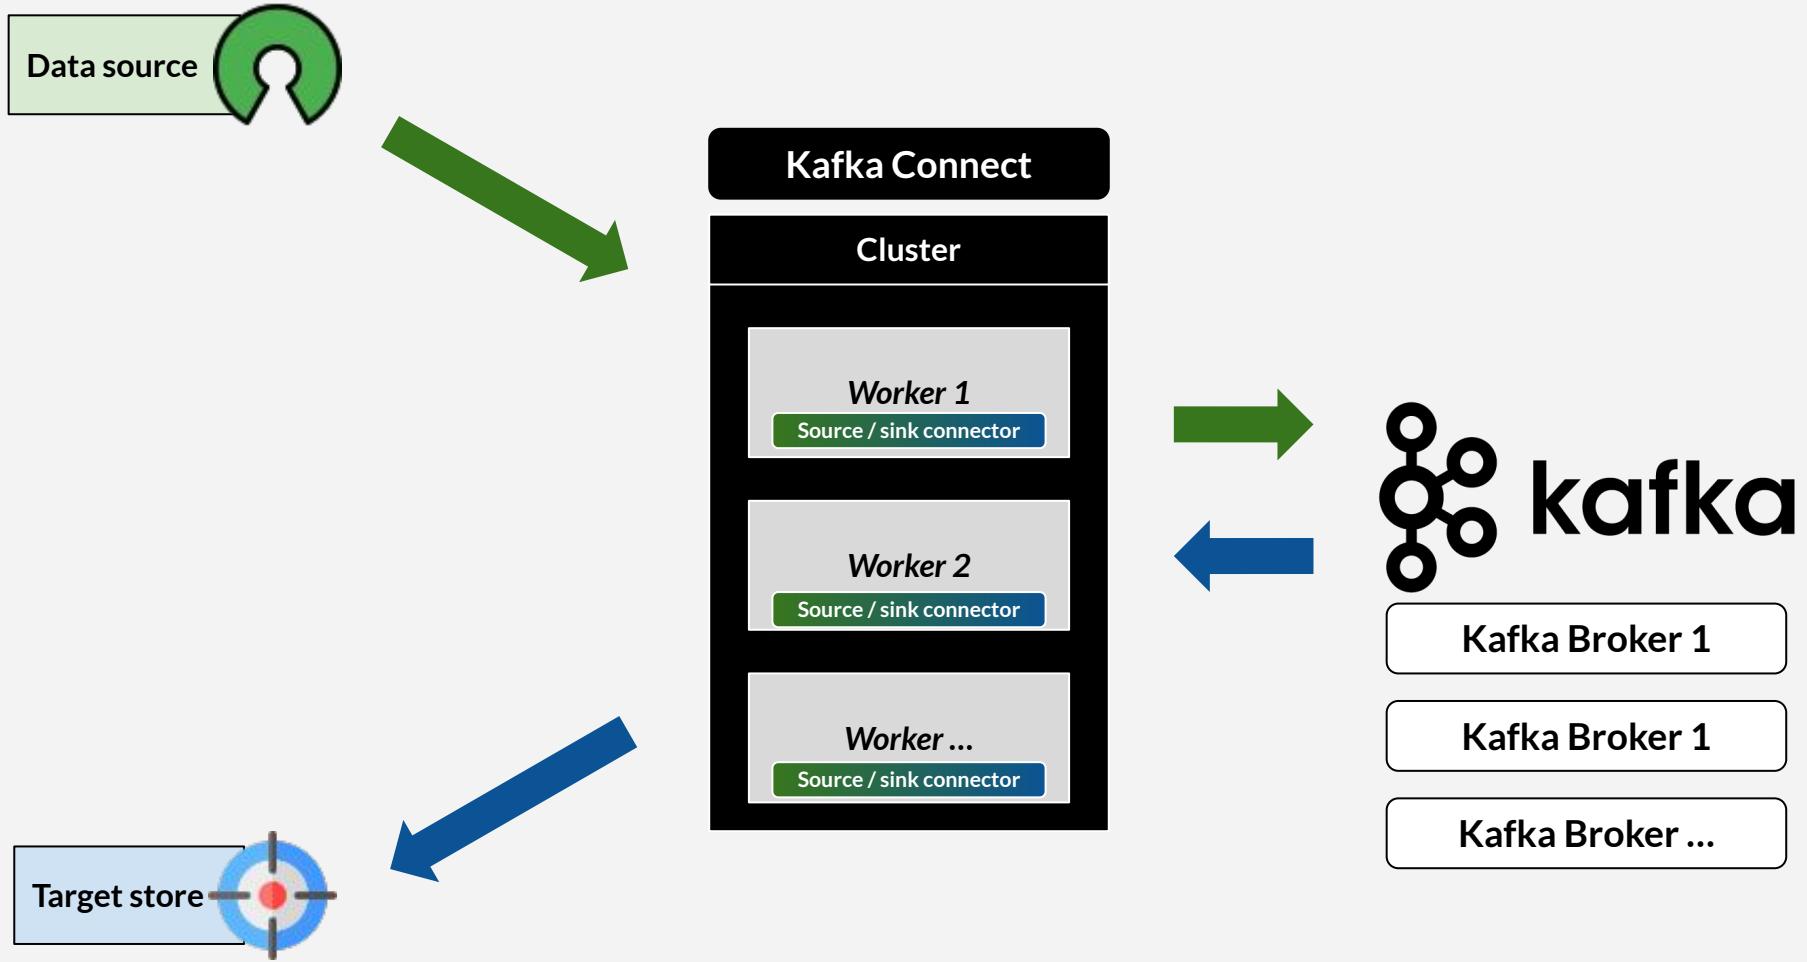
- × Write a lot of producers and consumers
- × Extra time & effort for performance & reliability
- × Good news : ***you don't have to write your own!***
- × People / companies already wrote them
- × Read plugin : from non-kafka into kafka
- × Write plugin : from kafka into non-kafka
- × **Kafka Connect**

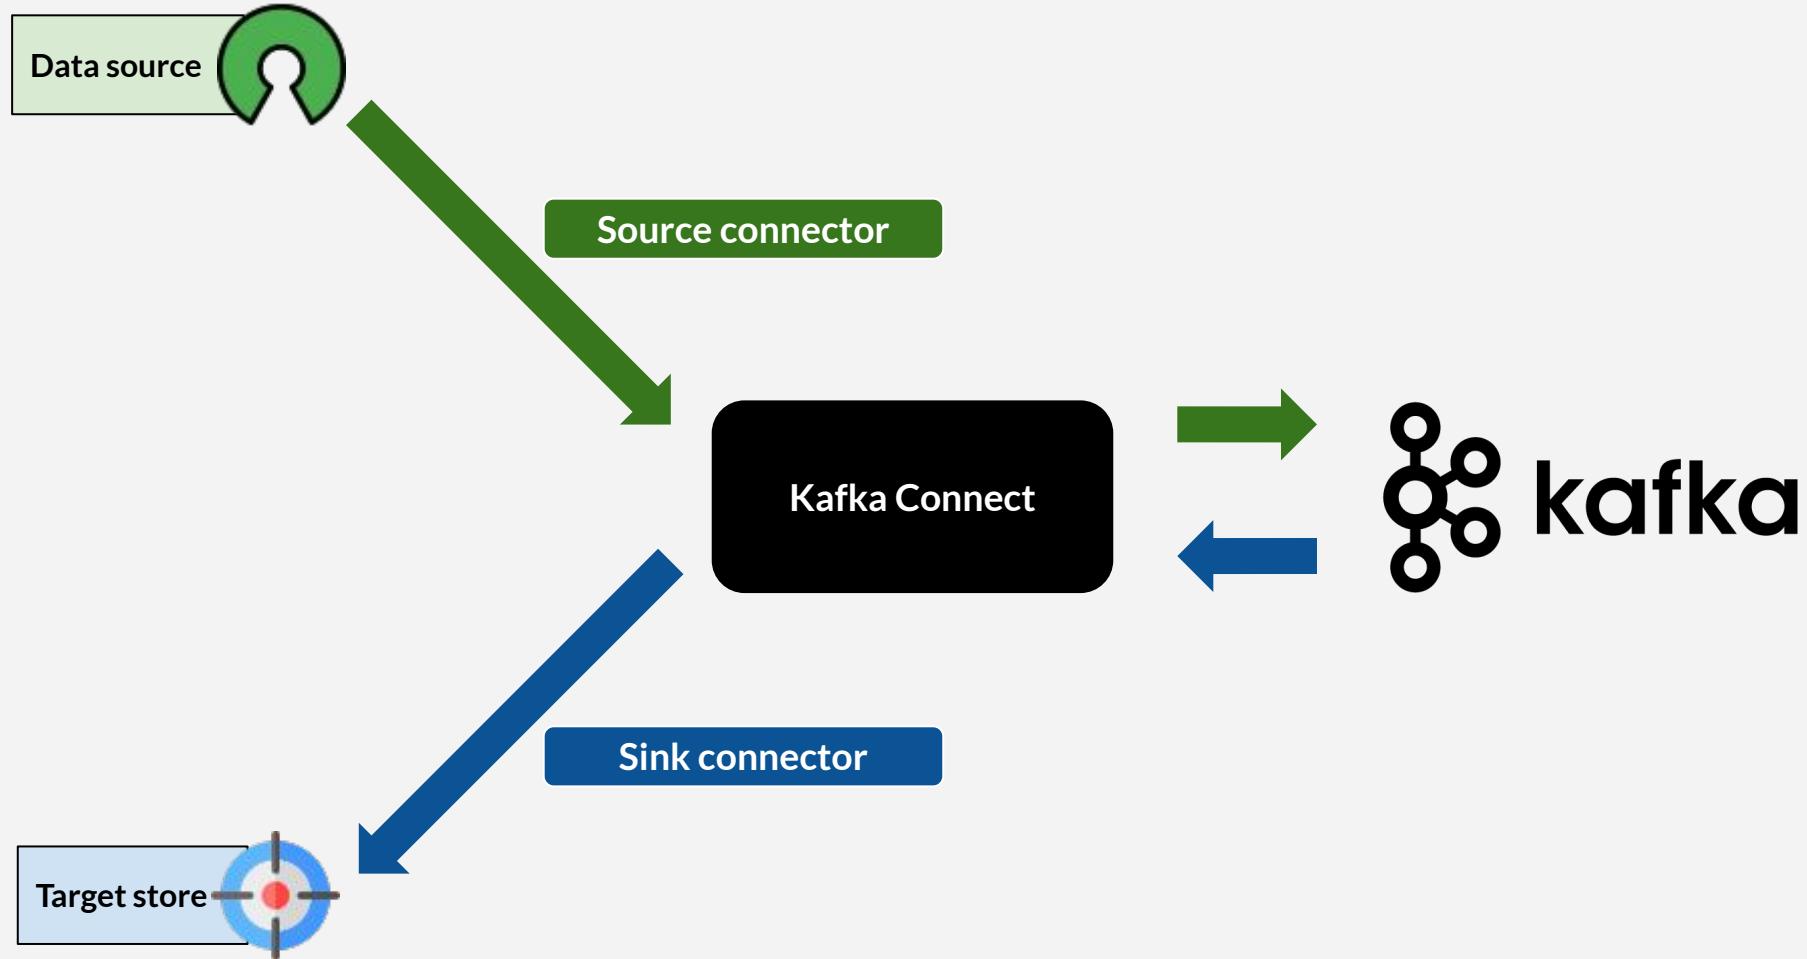


Introducing Kafka Connect









Kafka Connect

- × Additional platform for kafka
- × Data integration
- × Transfer data between Kafka - non kafka
- × Horizontally scalable & fault tolerant
- × Uses connectors for interact with kafka server



Connectors

- × Java jar file
- × Plugin for kafka connect
- × Interface between kafka and non-kafka
- × **Source** connector : read (ingest) into kafka (**producer**)
- × **Sink** connector : write from kafka to non-kafka (**consumer**)
- × Install connector for specific need
- × Write configuration (json)
- × Declarative configuration
- × Fast & reliable



Connectors

- × A lot of source / sink connectors
- × Shorten time and effort

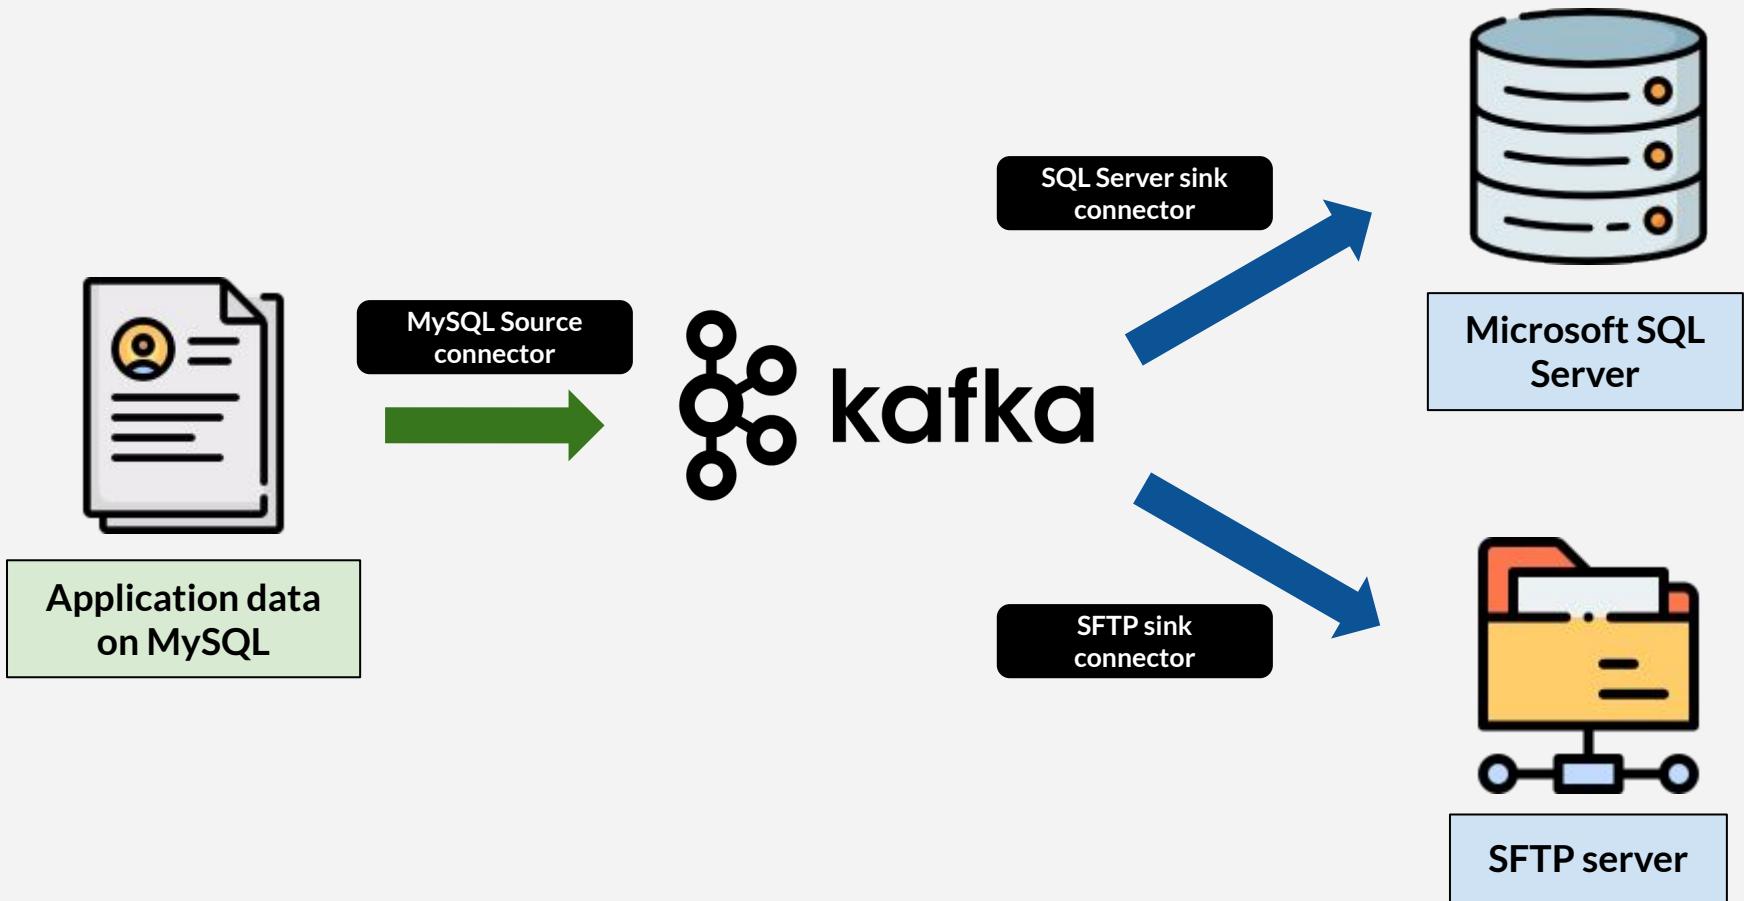


How to Get Connectors?

- × Curated list on [confluent.io/hub](https://www.confluent.io/hub)
- × Google : *kafka source / sink connector for xxx*
- × Build your own



Use Case : Write to Data Stores

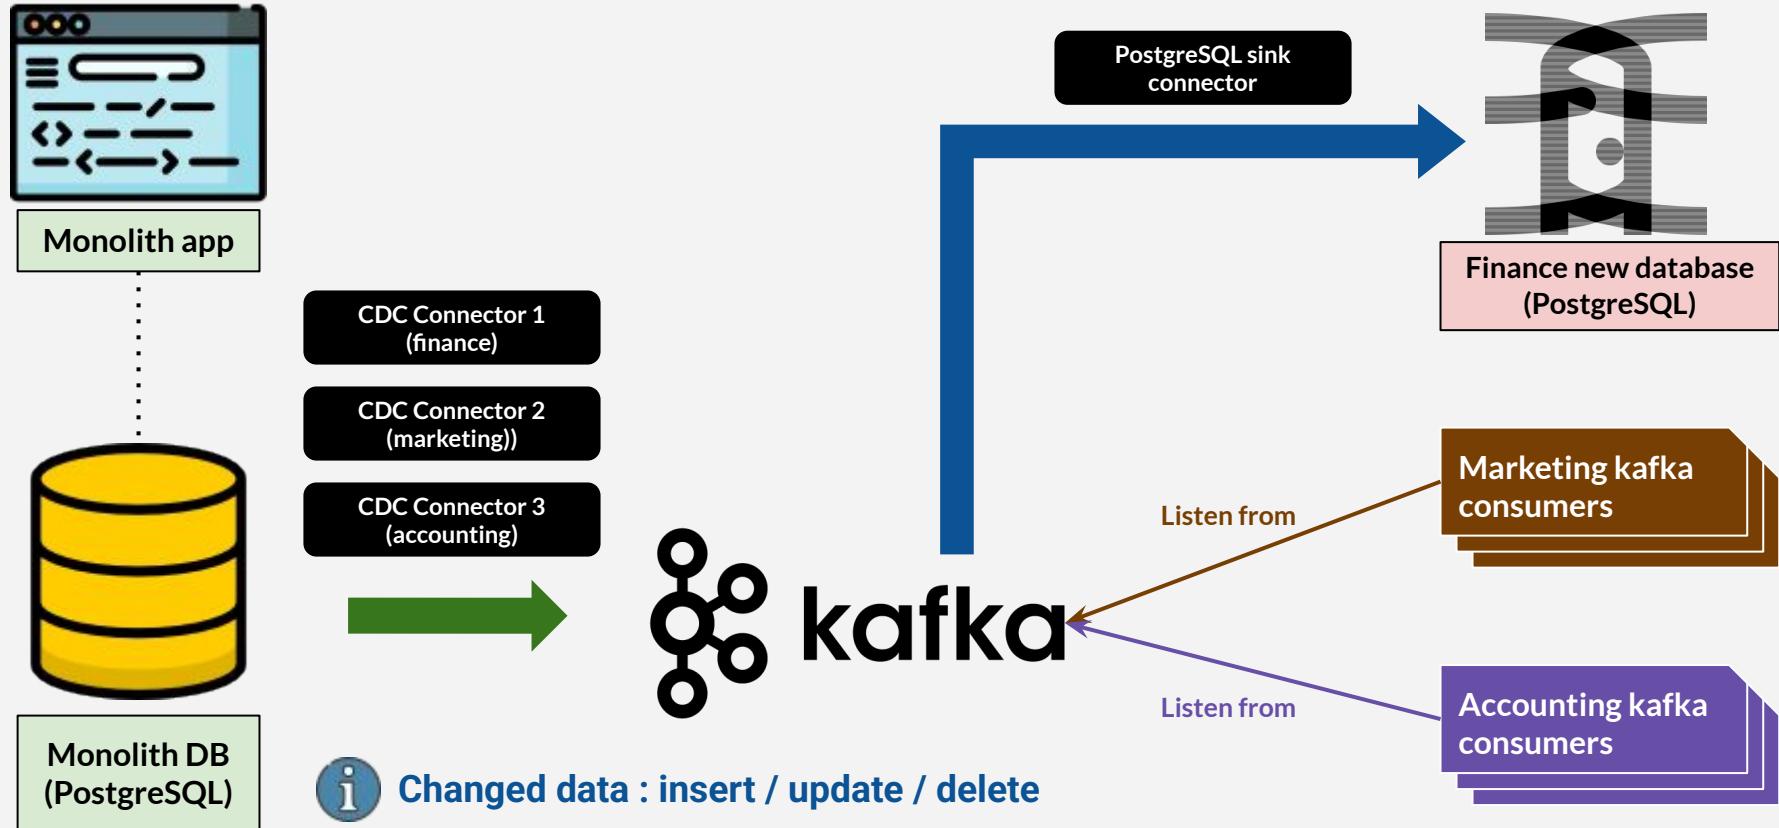


Use Case : Modernize Legacy System

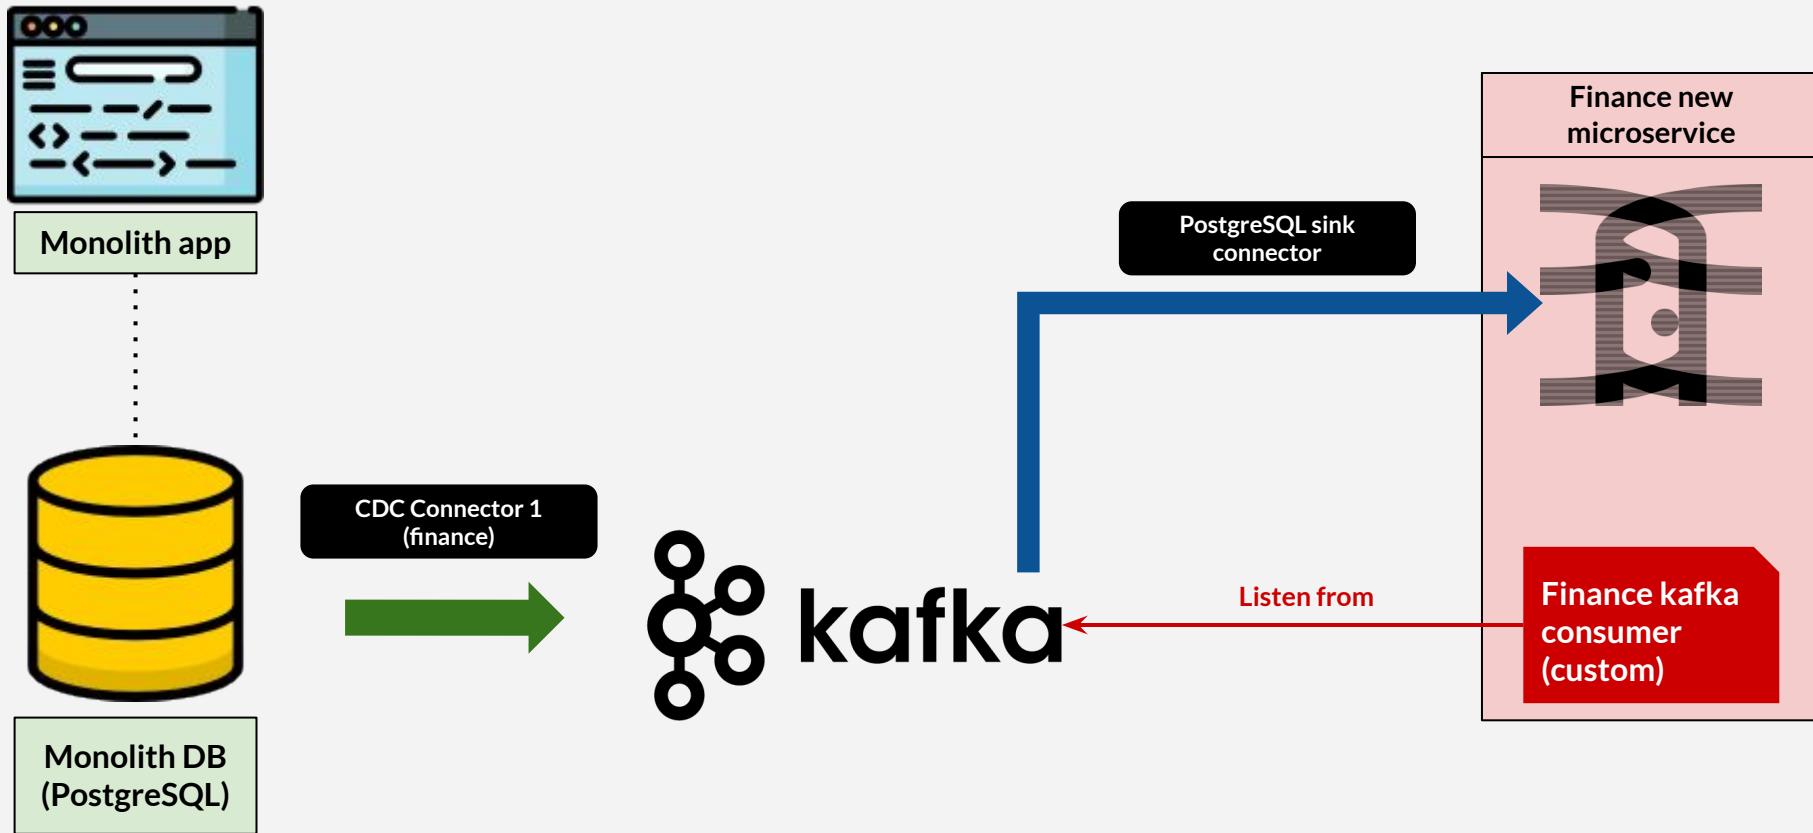
- × Modernize legacy monolith into microservices
- × Modernization is hard and long
- × Legacy system still needs to run during modernization
- × Modernize functionalities part by part
- × Use kafka connect CDC (Change Data Capture) connectors
- × *Microservice Architecture & Pattern course (and discount) available on last lecture of this course*



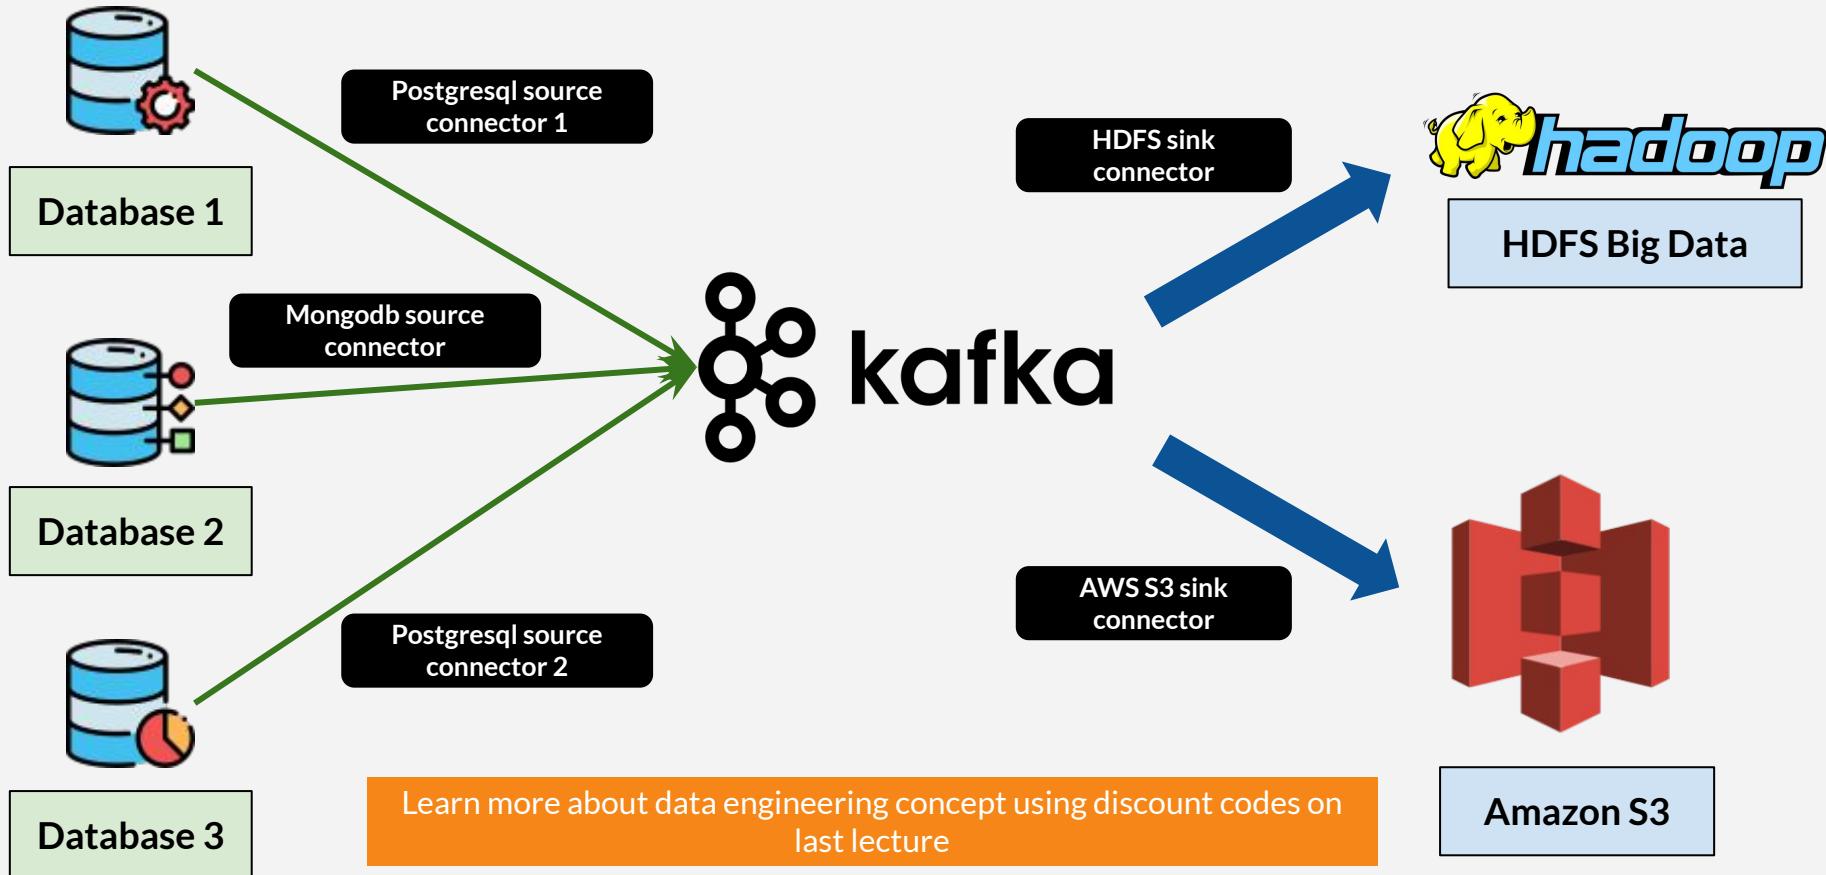
Use Case : Modernize Legacy System



Use Case : Modernize Legacy System



Use Case : Data Engineering ETL Pipeline



Kafka Connect

- × Cluster can has one or more workers (servers)
- × Connector + user configuration = **task**
- × A user configuration might spawn one or more task
- × Worker will execute task



Kafka Connect on Docker



```
#> docker-compose -f [script-file] -p [project] down
```

Part	What to run (in sequence)
1 - core kafka	<pre>#> docker-compose -f docker-compose-core.yml -p core up -d</pre>
2 - kafka connect	<pre>#> docker-compose -f docker-compose-core.yml -p core down #> docker-compose -f docker-compose-connect.yml -p connect up -d #> docker-compose -f docker-compose-connect-sample.yml -p connect-sample up -d</pre>
3 - kafka full	<pre>#> docker-compose -f docker-compose-connect.yml -p connect down #> docker-compose -f docker-compose-connect-sample.yml -p connect-sample down #> docker-compose -f docker-compose-full.yml -p full up -d #> docker-compose -f docker-compose-full-sample.yml -p full-sample up -d</pre>

Sample Use Cases



Use Cases

- × Sample use cases
 - × Basic connector usage
 - × Legacy modernization with CDC
 - × Data engineering
- × **Note**
 - × Need at least 4 GB free memory for docker
 - × Simple use cases only
 - × Data validity is not a concern
 - × Temporary data (stored on docker containers)



Sample Data

- × SFTP tool in this course : *Filezilla*
- × PostgreSQL tool in this course : *Dbeaver*
- × Feel free to use your own tools
- × Available on Resource & Reference
- × Sample data generated using mockaroo.com
- × Mockaroo.com schema available on Resource & Reference



Use Case : Basic Connector File Source





CSV source
connector



kafka

employee-* .csv

```
employee_id,first_name,last_name,email,gender,birth_date,salary
79-8035819,Hanny,Stubs,hstubs0@github.com,F,1996-02-02,7061
24-9676514,Val,Bonnick,vbonnick1@icq.com,M,1991-09-11,2657
61-3900911,Juanita,Lanfranconi,jlanfranconi2@rediff.com,F,1992-07-08,7241
28-0495338,Elsie,Brownlea,ebrownlea3@theglobeandmail.com,F,1985-11-03,8026
13-9980006,Chris,Kettlestringes,ckettlestringes4@indiegogo.com,M,1994-12-14,1015
```

```
{  
  "schema": {  
    "type": "data type",  
    "fields": {  
  
    }  
  },  
  "payload": {  
    "field1": "value 1",  
    "field2": "value 2"  
  }  
}
```

Embedded schema
Automatically generated by source connector
Required by most sink connector

Schema

- × Contract
- × Set of rules for message, to be obeyed
- × Field name, data type, mandatory / optional
- × Kafka connect able to work with dynamic data structure
- × Need to inform data structure (schema)
- × Kafka source connector auto generate schema



Use Case : Basic Connector Database Sink





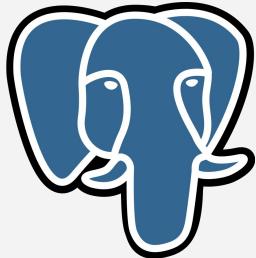
employee-* .csv

CSV source
connector



kafka

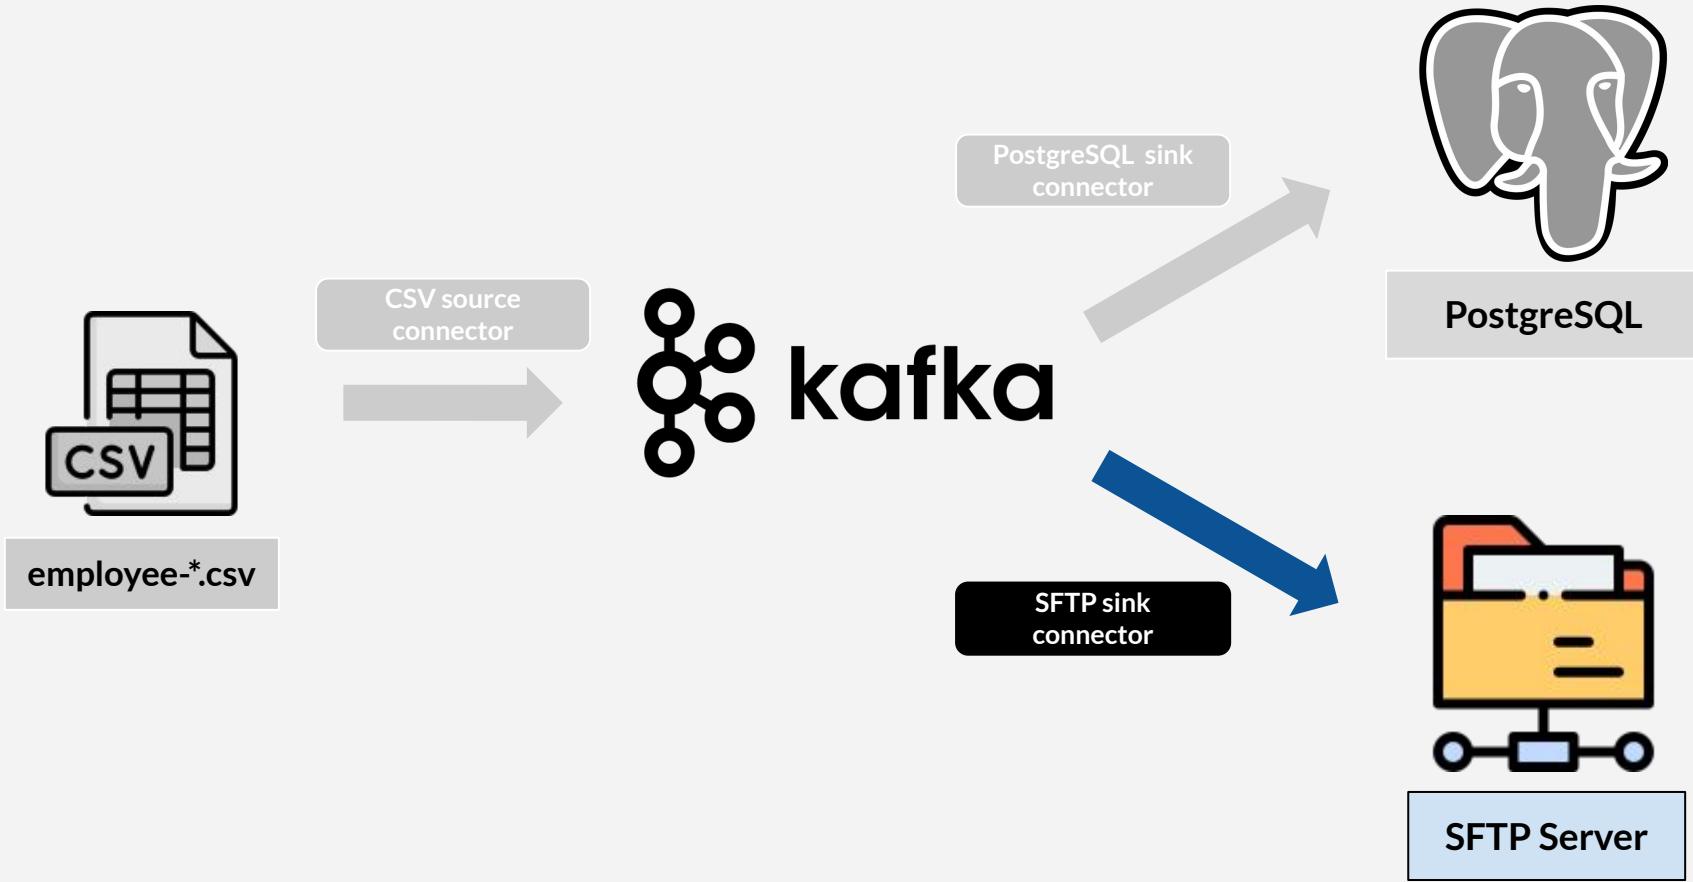
PostgreSQL sink
connector



PostgreSQL

Use Case : Basic Connector **SFTP Sink**





Use Case : Legacy Modernization Change Data Capture



Legacy Modernization

- ✗ Existing (legacy) application without kafka
- ✗ Rebuild into microservice with kafka
- ✗ Two systems (legacy & new) running together
- ✗ Synchronize data change (new, update, delete)
- ✗ Data source is from legacy application
- ✗ How?



Change Data Capture - 1st Way

- × Scheduler(hourly) that query data from legacy database based on *last_updated_date* field
- × But
 - × Not every table has *last_updated_date* field
 - × One hour is too long, need low latency (e.g. 60 seconds)
 - × Scheduler with short delay might problematic

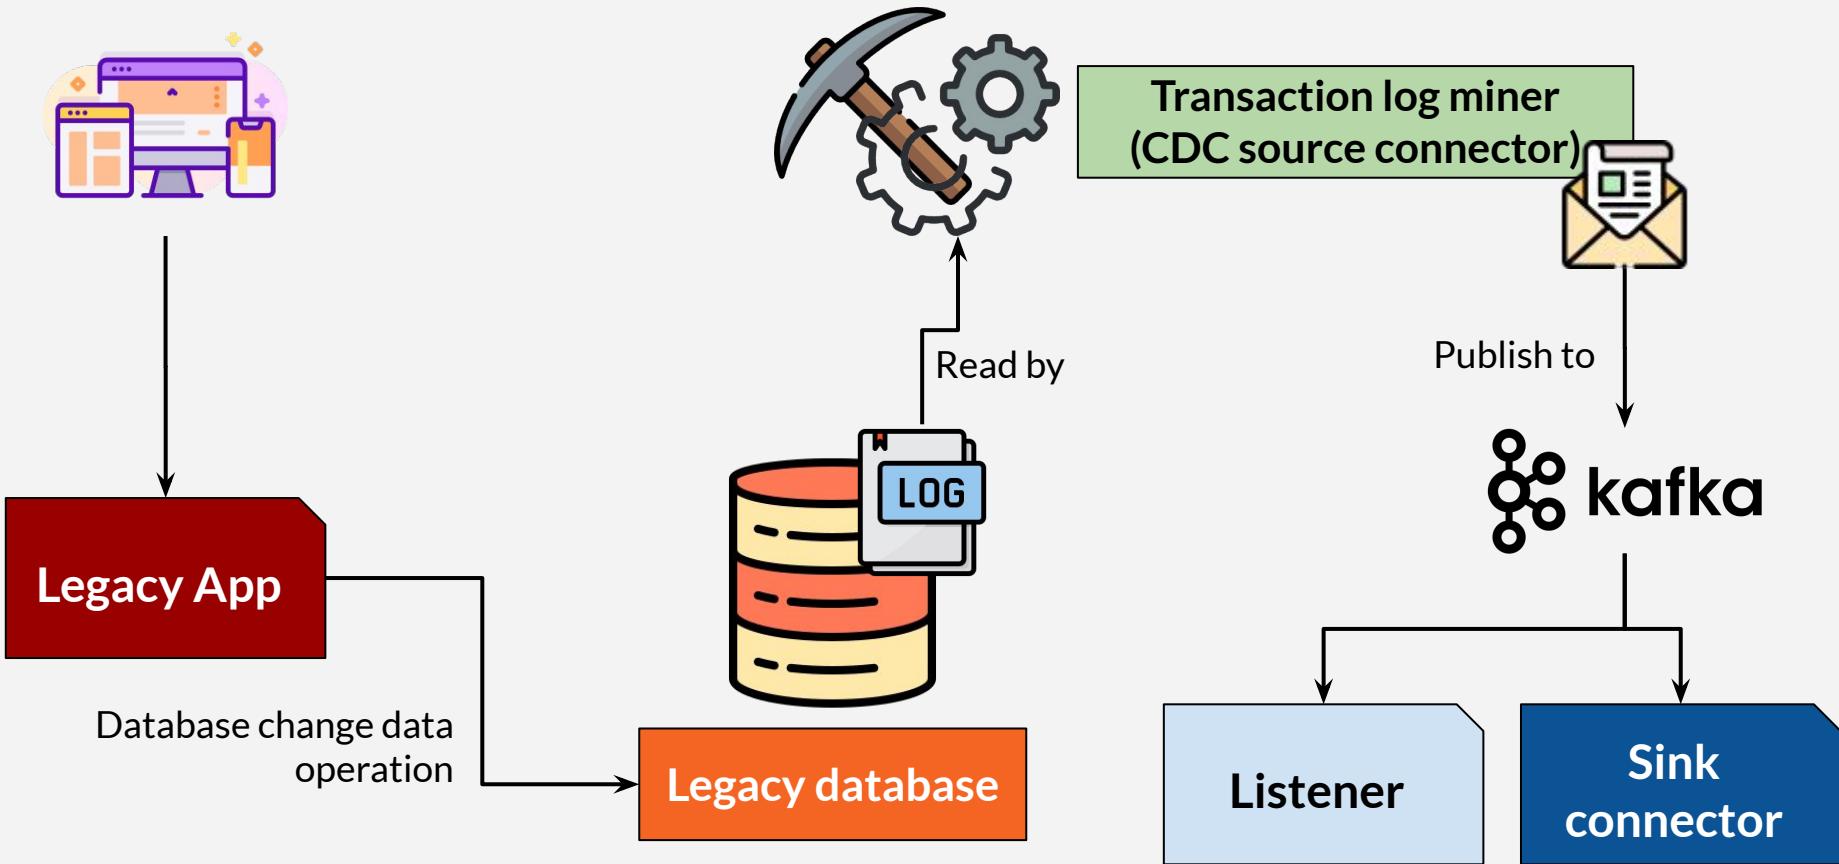


Change Data Capture - 2nd way

- × Each data change directly sent to kafka
- × Coding kafka integration points in legacy application is painful & not worth the effort
- × Good if we can implement this, but with minimal effort
- × Possible without writing code
- × Database transaction log tailing
- × Kafka connect with CDC (change data capture) source connector



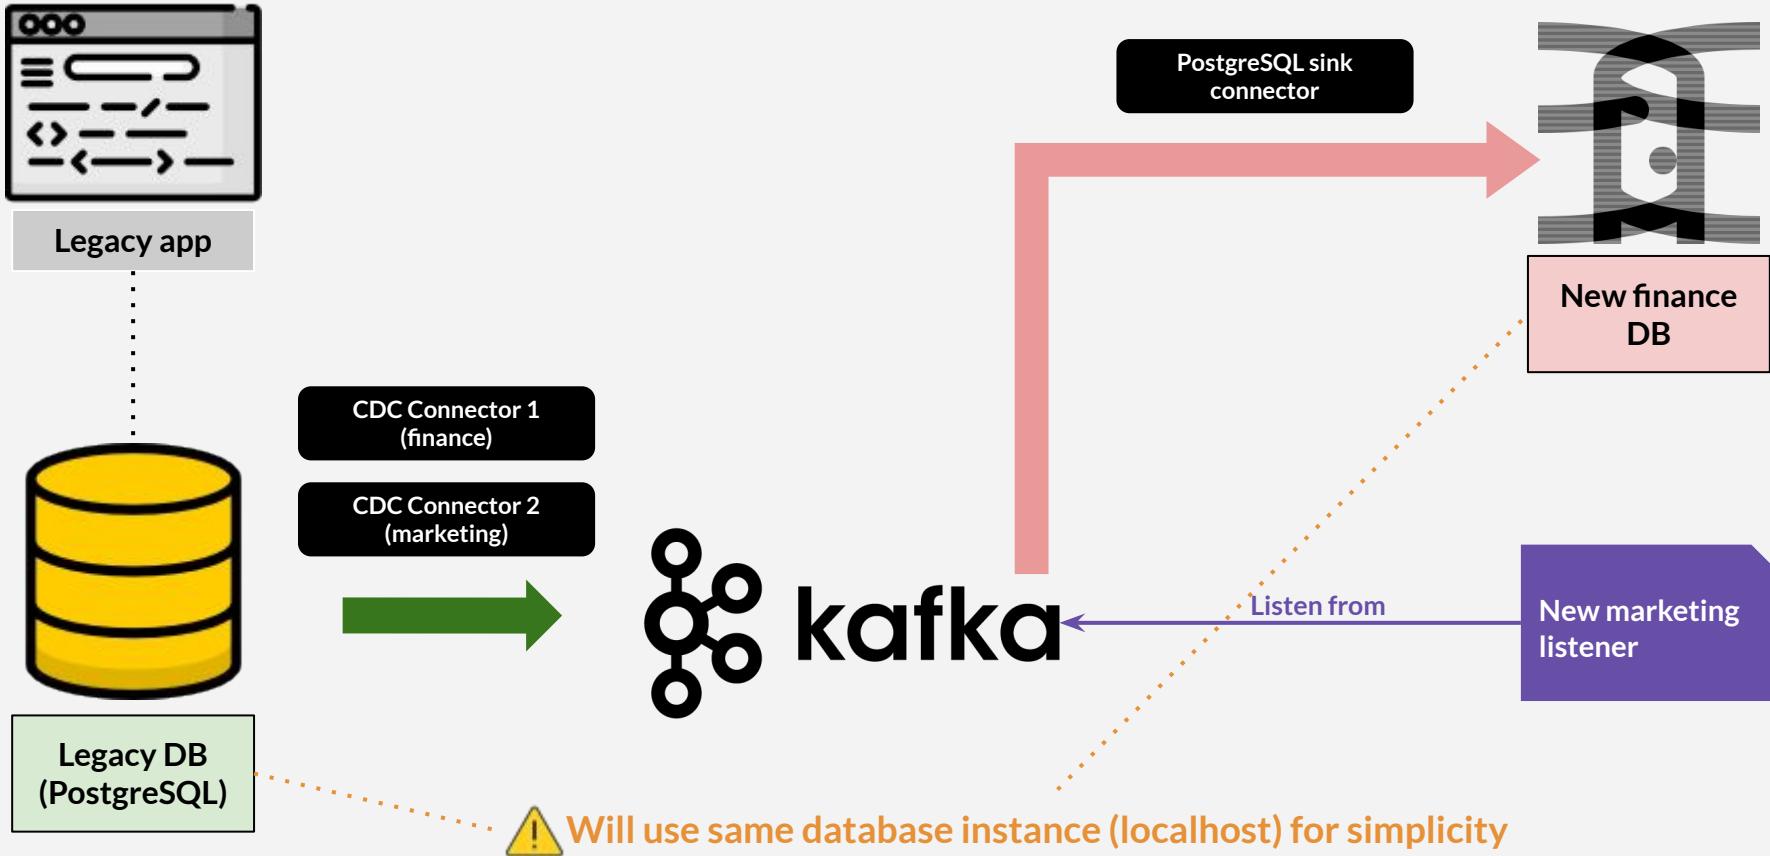
Transaction Log Tailing



Microservice Architecture & Pattern

- × Transaction log tailing is just one pattern to ease your life
- × Available in my course **Microservice Architecture & Pattern with Java & Kafka**
- × Grab discount code on last section of this course

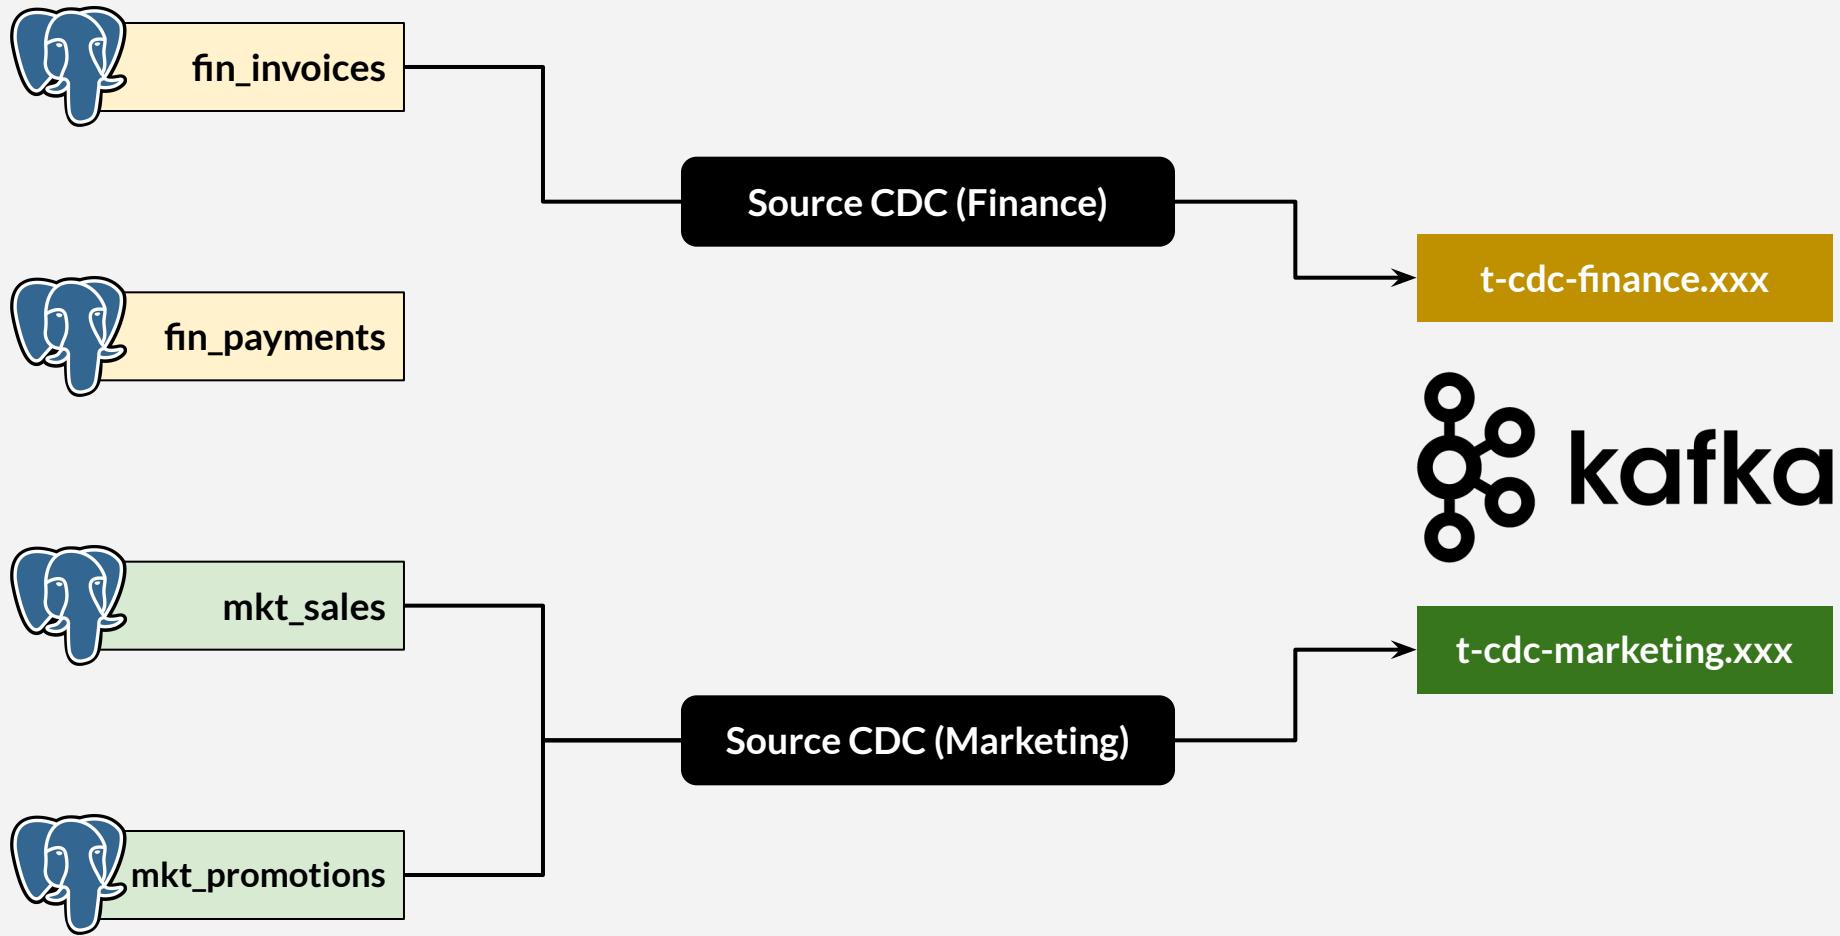




Use Case : Legacy Modernization

CDC PostgreSQL Connector





Use Case : Legacy Modernization

PostgreSQL Sink Connector



Use Case : Legacy Modernization Marketing Consumer



Spring Initializr

- × start.spring.io
- × Java project with gradle
- × Group : **com.course.kafka**
- × Artifact : **kafka-connect-sample**
- × Package name : **com.course.kafka**
- × Dependency : **Spring Kafka, Spring Kafka Stream**



Json Deserializer?

- × Why use default (String deserializer)
- × JSON (De)serializer is part of Spring
- × Spring specific mechanism to (de)serialize
- × Works out-of-the-box if publisher uses Spring JSON serializer
- × CDC connector is **not** part of spring framework
- × Process JSON string manually



It's not A One-Stop-Solution



Kafka Connect

- × Basic use case works
- × Complex use case needs more effort
- × In real life:
 - × Kafka connect can helps
 - × But it's not a magical problem solver
- × Need proper Kafka Connect configuration



Kafka Connect

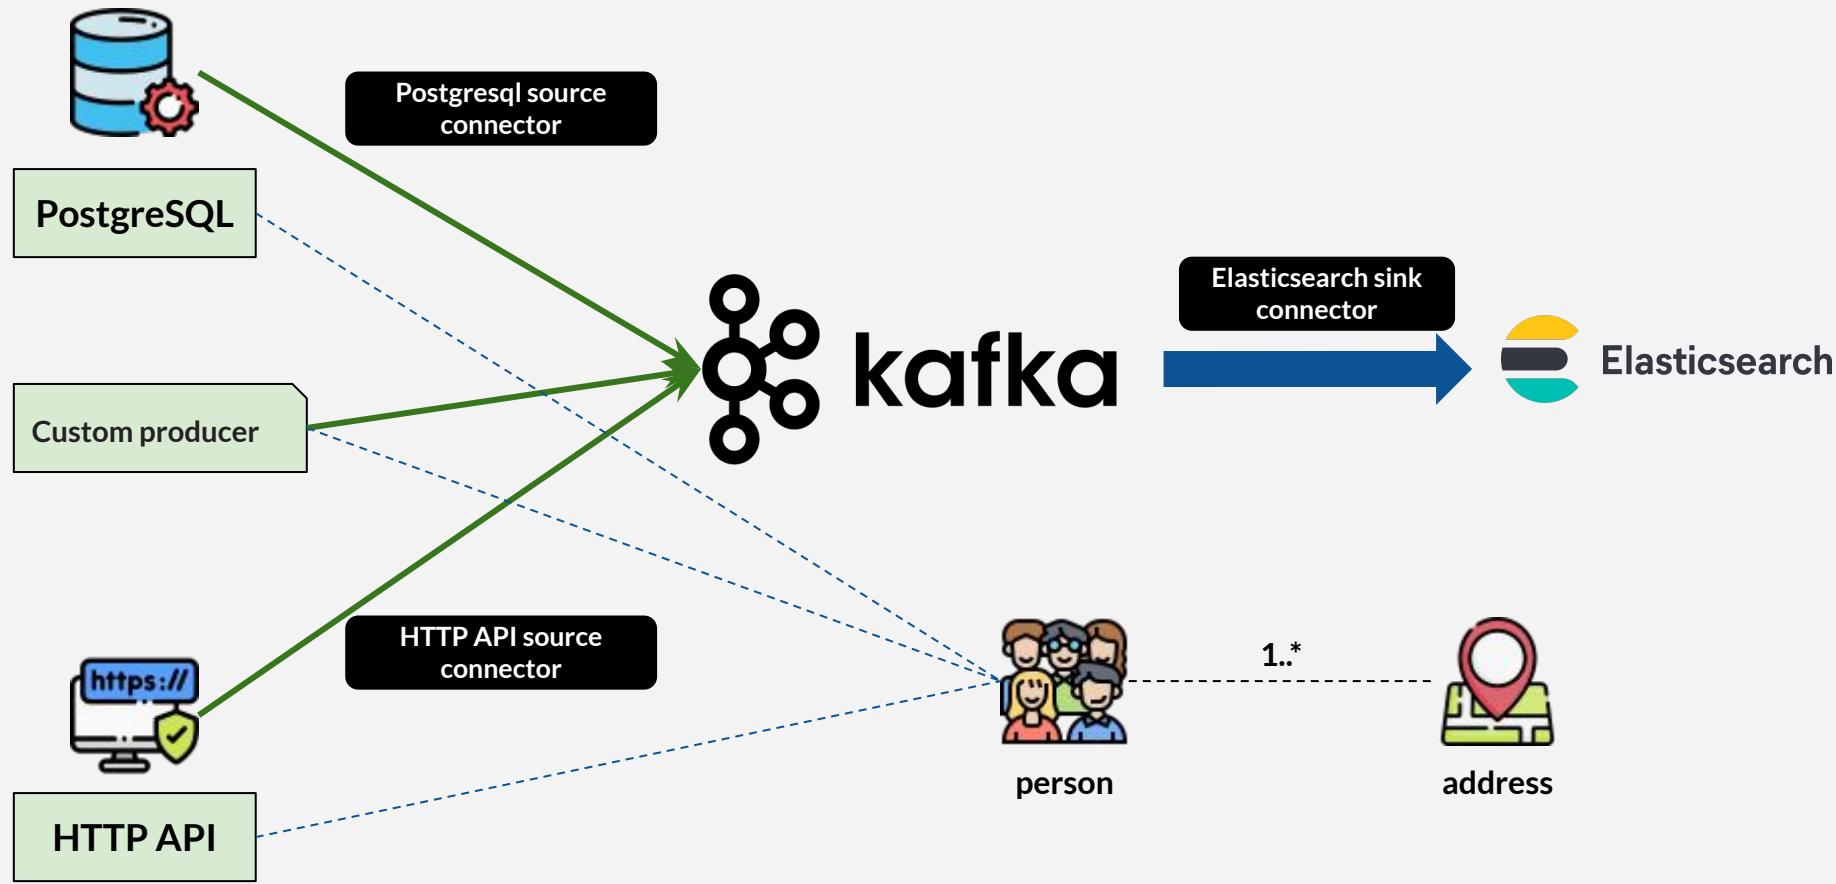
- ✗ Not entirely without cost
- ✗ Connector might need license
- ✗ Example : cloud storage
- ✗ Amazon / Google / Azure



Use Case : Data Engineering Database Source



Use Case : Data Engineering ETL Pipeline



Use Case : Data Engineering HTTP Source



Use Case : Data Engineering Custom Source



Data Format Differences



Database

```
{  
  "schema":{  
    ...  
  },  
  "payload":{  
    "person_id":2100,  
    "id_card_number":"401-99-8581",  
    "full_name":"Hamlin Kayes",  
    "email":"hkayes2r@bigcartel.com",  
    "address":"0 Westerfield Place",  
    "city":"Sājūr",  
    "postal_code":null  
  }  
}
```

*Schema + payload
Actual data in payload
Using snake_case
Single address*

HTTP

```
{  
  "schema":{  
    ...  
  },  
  "payload":{  
    "value":{  
      "id_card_number":"444-25-9069",  
      "full_name":"Rebekkah Toller",  
      "email":"rtoller8@ezinearticles.com",  
      "addresses":{  
        "address_id":6710,  
        "address":"5 Lillian Parkway",  
        "city":"Lyubim",  
        "postal_code":152470},  
        "address_id":7860,  
        "address":86278 Calypso Trail,  
        "city":Gaotieling,  
        "postal_code":null  
      },  
      "key":9d38e4bf-3b1d-32d9-b78b-7a8135233597,  
      "timestamp":1644280143619  
    }  
  }  
}
```

*Schema + payload
Actual data in payload.value
Using snake_case
Array address*

Custom

```
{  
  "personId":84527,  
  "fullName":"Joe Auer",  
  "email":"joeauer99@hotmail.com",  
  "addresses": [  
    {  
      "addressId":64288,  
      "address":"4639 Everett Walks",  
      "city":"Stacyfort",  
      "postalCode":null  
    }  
  ]  
}
```

*No schema, no payload
Represents actual data
Using camelCase
Array address*

Database

```
{  
  "schema": {  
    ...  
  },  
  "payload": 3102  
}
```

HTTP

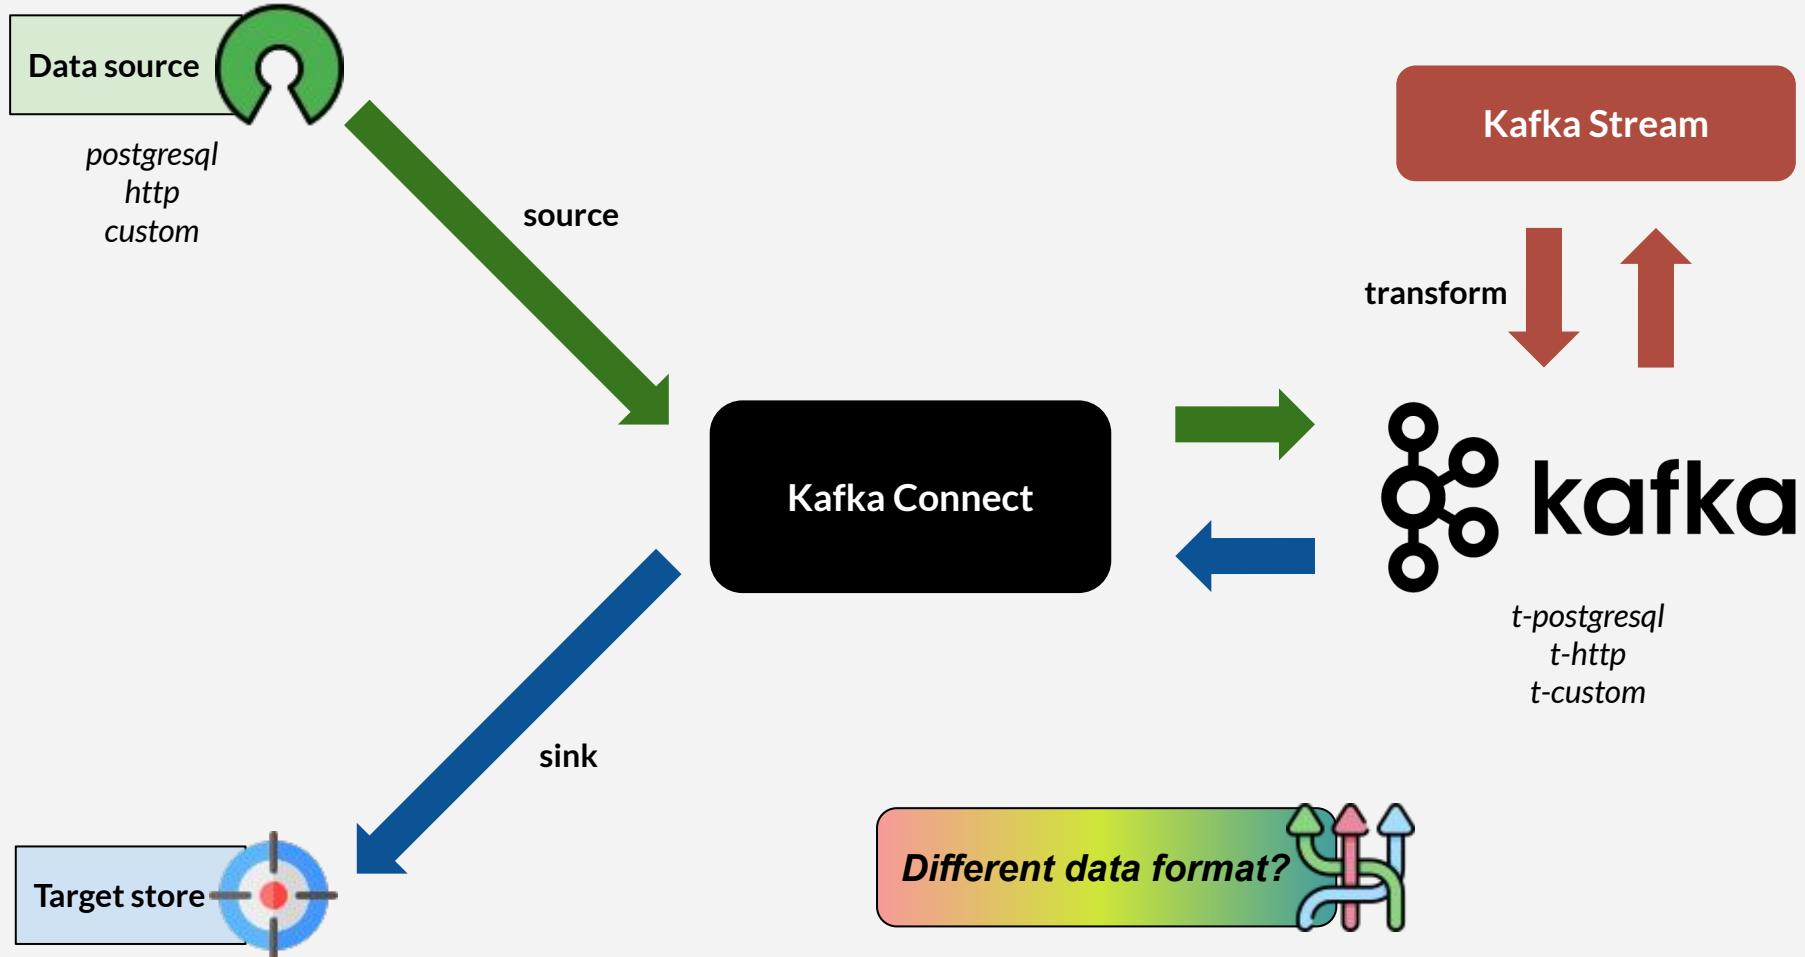
```
{  
  "schema": {  
    ...  
  },  
  "payload": {  
    "key": "62aeb383-539a-3ea5-99cb-a16fb4b3681b"  
  }  
}
```

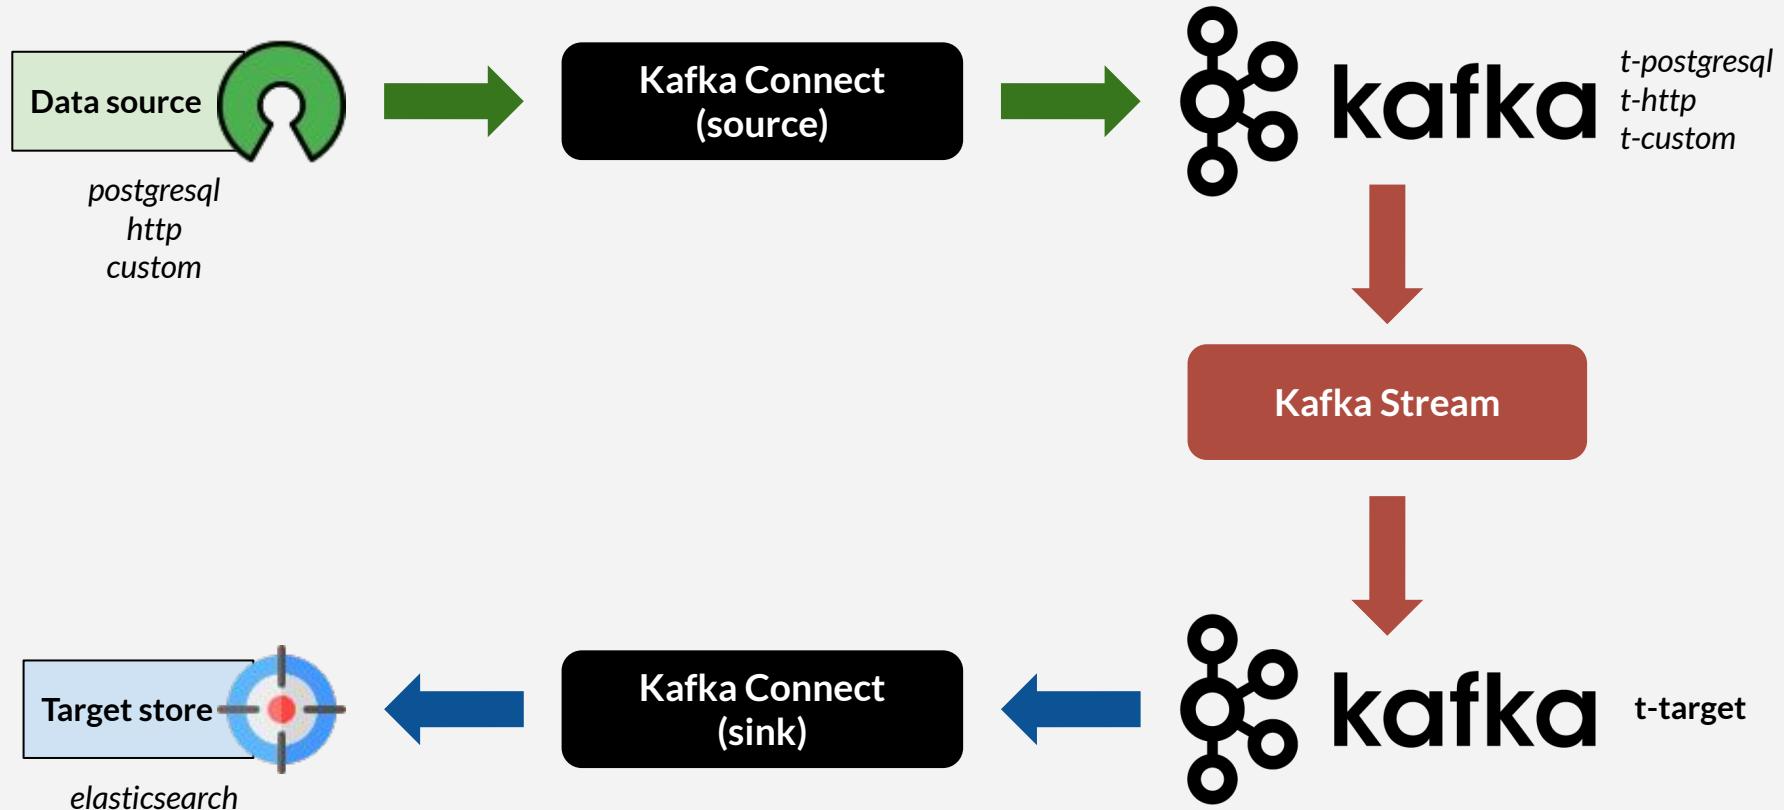
*Schema + payload
Actual data in payload.key
Random UUID*

Custom

*Schema + payload
Actual data in payload
Address ID*

No key





Kafka Stream & Kafka Connect Code Overview

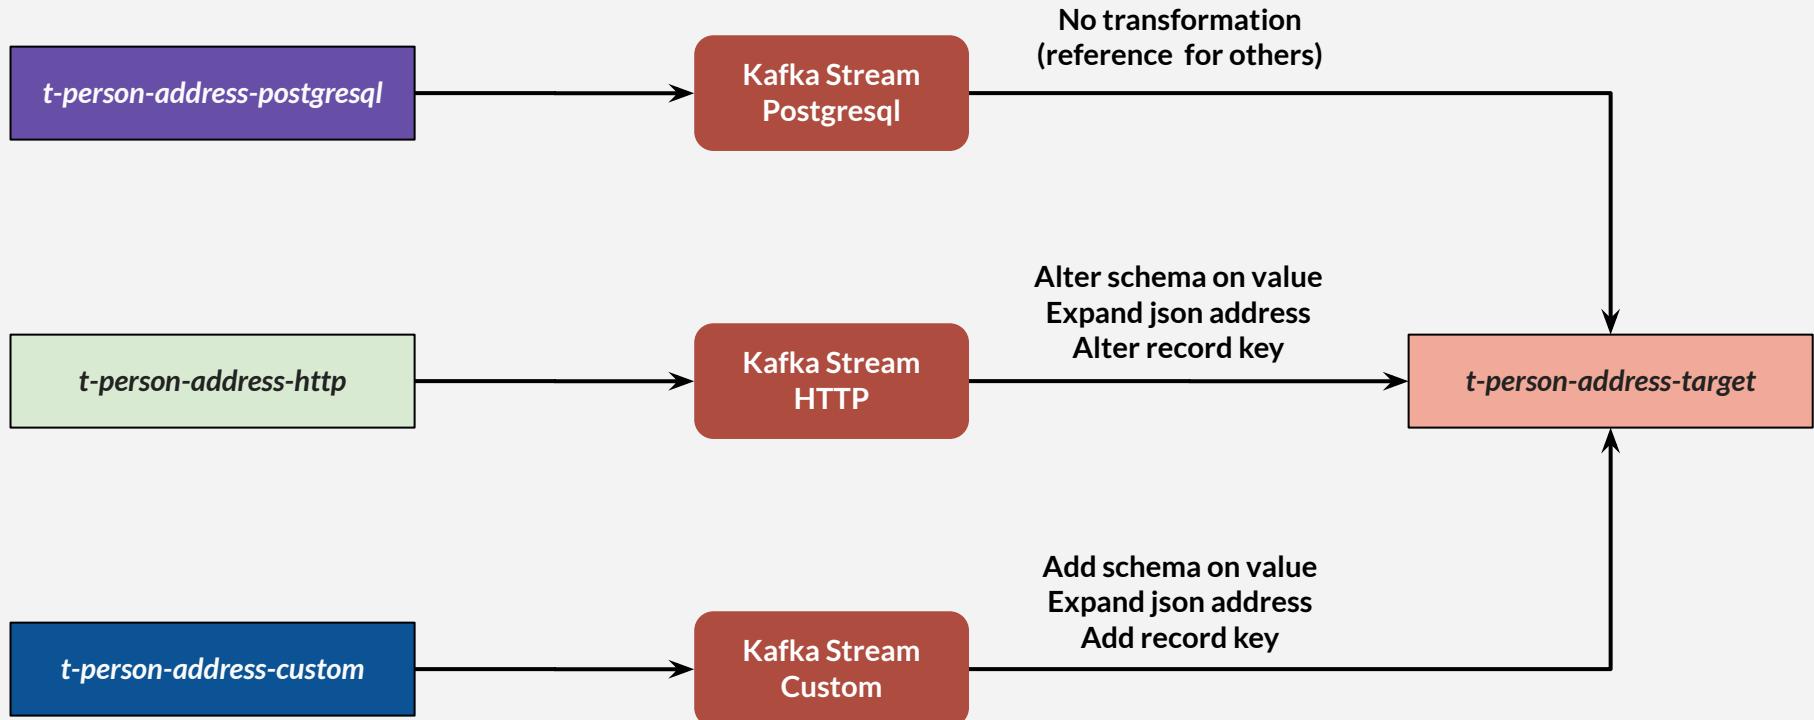


Kafka Stream & Kafka Connect

- × Source topics (from source connectors)
 - × *t-person-address-postgresql*
 - × *t-person-address-http*
 - × *t-person-address-custom*
- × Target topic (for sink connector)
 - × *t-person-address-target*
- × Use java generic coding practice
- × String serde for key and value



Kafka Stream Transformation



All

<i>Topic</i>	<i>Class</i>	<i>Description</i>
all	message.KafkaConnectMessage	Wrapper class for kafka connect message that contains schema and payload . Payload uses java generic.
	schema.KafkaConnectSchema	Represents schema on KafkaConnectMessage

PostgreSQL Source

Topic	Class	Description
<i>t-person-address-postgresql</i>	-	No need to define class representing message. This is the message format we will use on target topic.
	<code>stream.PersonAddressFromPostgresqlStream</code>	<p>Kafka stream:</p> <ol style="list-style-type: none">1. Take data from <i>t-person-address-postgresql</i>2. Send as is, to <i>t-person-address-target</i> <p>No transformation needed, since the record format (key & value) is what we expect.</p>

HTTP Source

Topic	Class	Description
<i>t-person-address-http</i>	<code>message.KafkaConnectPersonAddressFromHttpMessage</code>	Represents message retrieved by HTTP source connector (contains full json string on <code>payload.value</code>)
	<code>message.KafkaConnectPersonMessageSnakeCase</code> <code>message.KafkaConnectAddressMessageSnakeCase</code>	Holder for converting json string in <code>payload.value</code> into java object (the actual person & address data). Fields are <code>snake_case</code> .
	<code>stream.PersonAddressFromHttpStream</code>	Kafka stream: <ol style="list-style-type: none">1. Take data from <i>t-person-address-http</i>2. Transform (expand json <code>addresses</code> and convert each address into one record). Alter schema on record key & value3. Send to <i>t-person-address-target</i>

Custom Source

Topic	Class	Description
<i>t-person-address-custom</i>	<code>message.PersonMessage</code> <code>message.PersonAddress</code>	Represents message published by dummy scheduler. Fields are camelCase.
	<code>stream.PersonAddressFromCustomStream</code>	Kafka stream: <ol style="list-style-type: none">1. Take data from <i>t-person-address-custom</i>2. Transform (expand json addresses and convert each address into one record). Add schema on record key & value3. Send to <i>t-person-address-target</i>

Target Sink (From All Sources)

Topic	Class	Description
<i>t-person-address-target</i>	<code>message.KafkaConnectPersonTargetMessage</code>	<p>Represents payload on record value (class <code>KafkaConnectMessage</code>) to be published to <i>t-person-address-target</i>. This actually the same format as data published by JDBC source connector at <i>t-person-address-postgresql</i></p>
	<code>schema.KafkaConnectPersonAddressTargetKeySchema</code>	<p>Singleton for creating schema on record key to be published to <i>t-person-address-target</i>.</p>
	<code>schema.KafkaConnectPersonAddressTargetValueSchema</code>	<p>Singleton for creating schema on record value to be published to <i>t-person-address-target</i>.</p>

Kafka Stream & Kafka Connect Code Hands on



Tip : Override Converter



Use Case : Data Engineering Elasticsearch Sink



AI Assistant and Kafka Connect



AI Assistant & Kafka Connect

- × Can use AI assistant up to a certain point
- × Generate basic Kafka Connect JSON configuration
- × Sometimes fine-tuning by human is better
 - × Limited / outdated AI assistant knowledge base
 - × Might need long & detailed instruction
 - × Might be faster to read & configure the connector based on recent documentation



AI Assistant & Kafka Connect

- ✗ Not much AI assistant usage in the Kafka Connect lessons
- ✗ Some prompts example provided on **Resources & References**
- ✗ Generate JSON body that resembles connectors in this course
- ✗ Only some examples
- ✗ Difference between AI assistant vs manual
- ✗ AI assistant output is non-deterministic
- ✗ AI assistant generated JSON body will not be executed in this course



AI Assistant & Kafka Connect

- × Explain existing configuration
- × Ask AI assistant to explain existing JSON configuration



Kafka User Interface



Kafka GUI

- ✗ Use GUI instead console
- ✗ Available on market
- ✗ Might need license
- ✗ Check tool's website for pricing & license



Kafka GUI

- × Confluent Control Center (confluent.io)
- × Kafdrop (github.com/obsidiandynamics/kafdrop)
- × Kafka-ui (github.com/provectus/kafka-ui)
- × Lenses (lenses.io)
- × Conduktor (conduktor.io)
 - × We will use this



Conduktor



Conduktor

- × Using new topics
- × Dummy producer, consumer
- × Dummy kafka stream
- × Source code available on **Resource & Reference**
- × *Note : Conduktor features might different from the course*

